



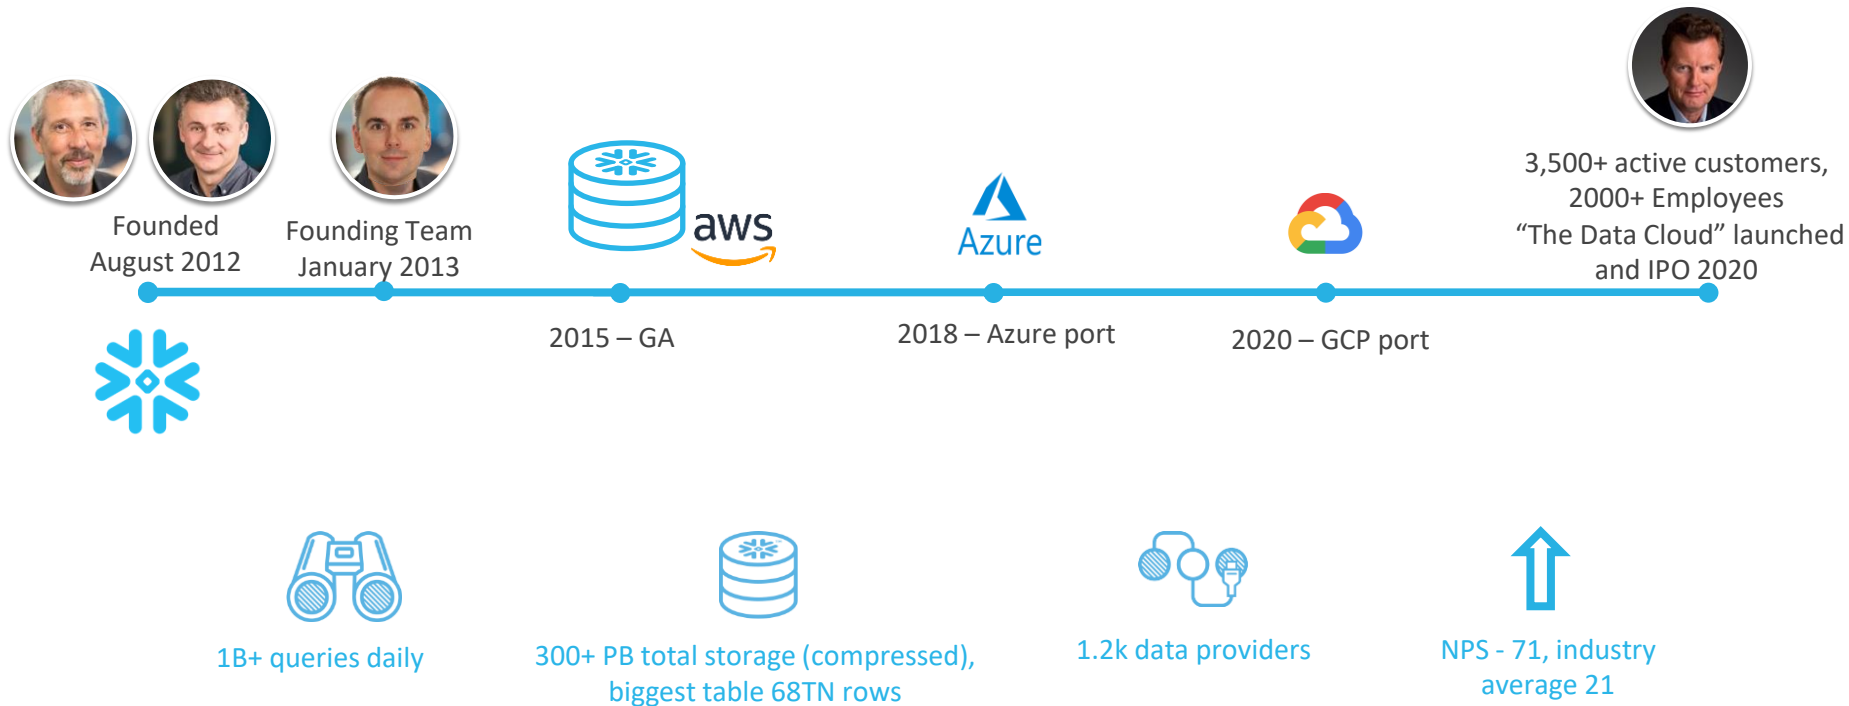
SNOWFLAKE DATA CLOUD

SHENZHEN ARCHSUMMIT 2021

Jiaqi Yan

Principal Software Engineer

QUICK HISTORY



WHY SNOWFLAKE?

No Good Solution to Tackle Modern Data Challenges

More Data



More Users



Faster Answers



Diverse Data Sources



SNOWFLAKE DATA CLOUD

Data Analytics Platform Built for the Cloud Era

UNLIMITED SCALE

- All Data
- All Users
- Instant Elasticity
- One System
 - Multi-region
 - Multi-cloud
- No Compromise

COLLABORATIVE

- Data Sharing
- Data Services
- Existing Content

SIMPLICITY

- Self-Managed Service
- No Tuning Knobs
- Democratize Data Analytics




UNLIMITED SCALE



TRADITIONAL DATABASE ARCHITECTURES


Limited Scalability, Not Elastic

Shared-nothing



- Distributed Storage
- Single Cluster
- Adopted by Gamma, Teradata, Redshift, Vertica, Netezza, ...

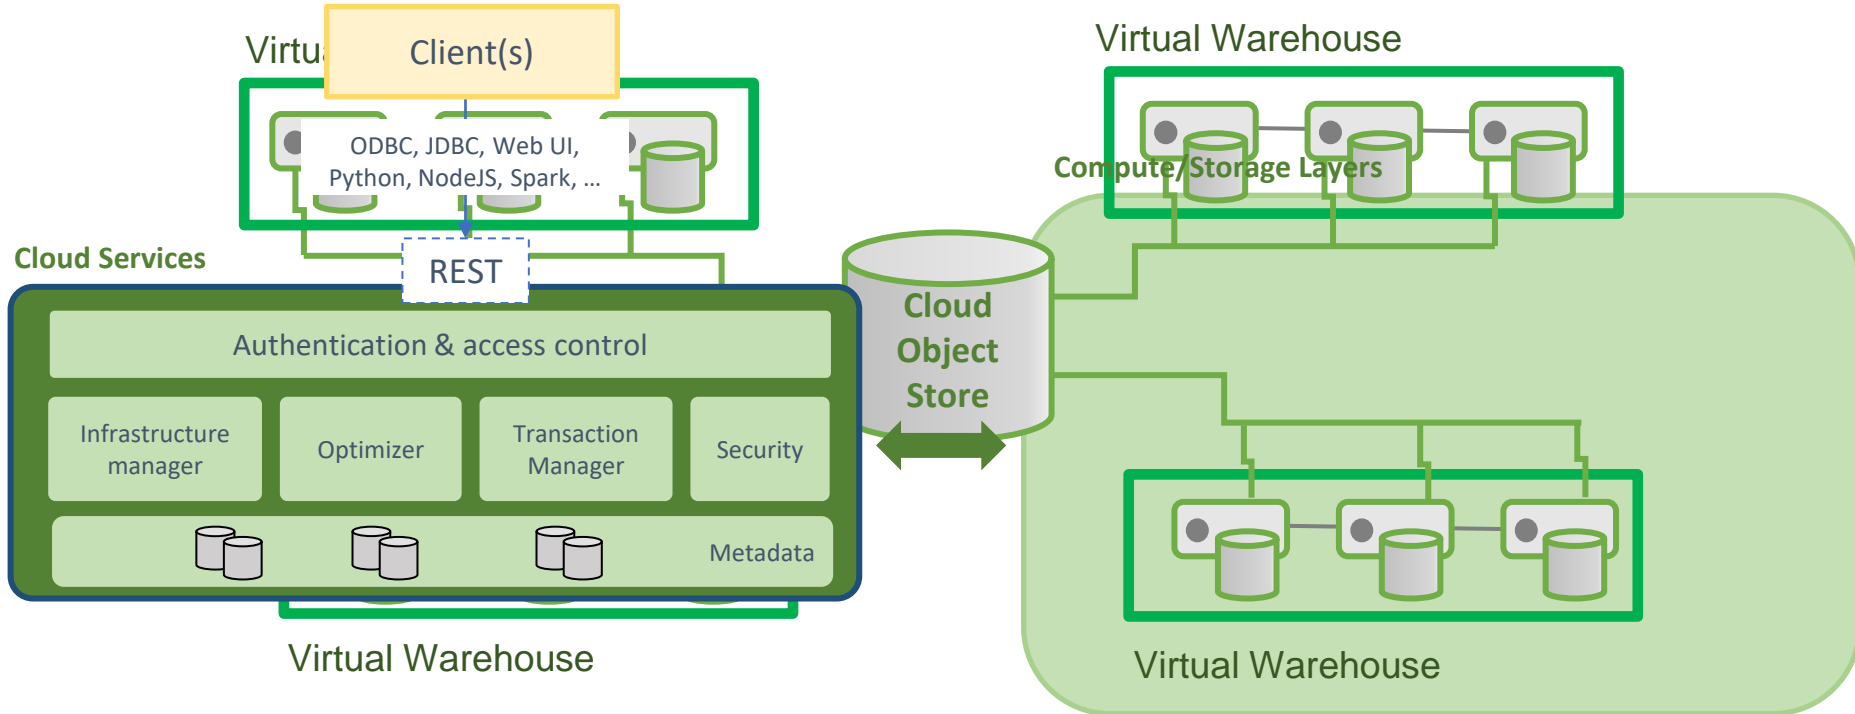
Shared-disk



- Centralized Storage
- Single Cluster
- Adopted by Oracle, Hadoop

SNOWFLAKE REGION ARCHITECTURE

Multi-cluster, Shared-data




STORAGE TIER



- **Immutable Storage**
 - Each table is automatically partitioned horizontally
 - Partition size is kept very small, generally 16MB
 - Each partition is backed by an immutable file
 - Partitions are columnar organized, compressed, encrypted
 - Partitions are the unit of change for transactions
- **Semi-structured**
 - Variant data type used to store schemaless semi-structured data
 - Automatic columnarization of semi-structured attributes
- **Partition Metadata**
 - Out-of-box, metadata is automatically stored for all columns/sub-columns in a partition
 - Leverage that metadata to perform partition pruning
 - Re-clustering service to improve pruning
 - Track all table mutations to provide full ACID support

COMPUTE TIER



- **Virtual warehouse**

- Snowflake Entity used to manage the set of compute resources used by a workload
- Made of one or more compute clusters
- Cluster size range from one to several hundred nodes
- Workloads are fully isolated from each other


- **Just-in-time Compute**

- Sub-second auto-resume when associated workload starts
- Online resize up and down possible while workload runs
- Auto-suspend when workload is no longer running
- Snowflake charges usage by second of compute resource used
➔ **FAST is free!**

- **Partition Cache**

- Node local memory and SSD storage used to cache partitions
- Only columns/sub-columns which are accessed are cached
- Highly available, fully stateless

CLOUD SERVICES



- **Control Plane of a Snowflake Region**

- Connection Management
- Infrastructure Provisioning and Management
- Metadata storage (use FDB) & management
- Query planning and optimization
- Transaction management
- Security management

- **Self-managed**


- Self-upgrade of both software and hardware
- Self-healing: replacement of failed servers and transparent re-execution of any failed queries
- Highly available over multiple availability zone
- Stateless: persistent sessions for load-balancing and transparent fail-over

SNOWFLAKE DATA CLOUD

One Integrated Platform Supporting Multiple Workload Types



Complete SQL
ACID
Low-latency
High-concurrency
UDFs, UDTs
Data Governance
Stored Procedures



Streaming Ingest
Tasks
Table Streams
External Functions
Data Pipelines



Semi-structured Data
Unstructured Data
External Tables




Java/Scala/Python
Data Frames



Rest APIs
Real-time

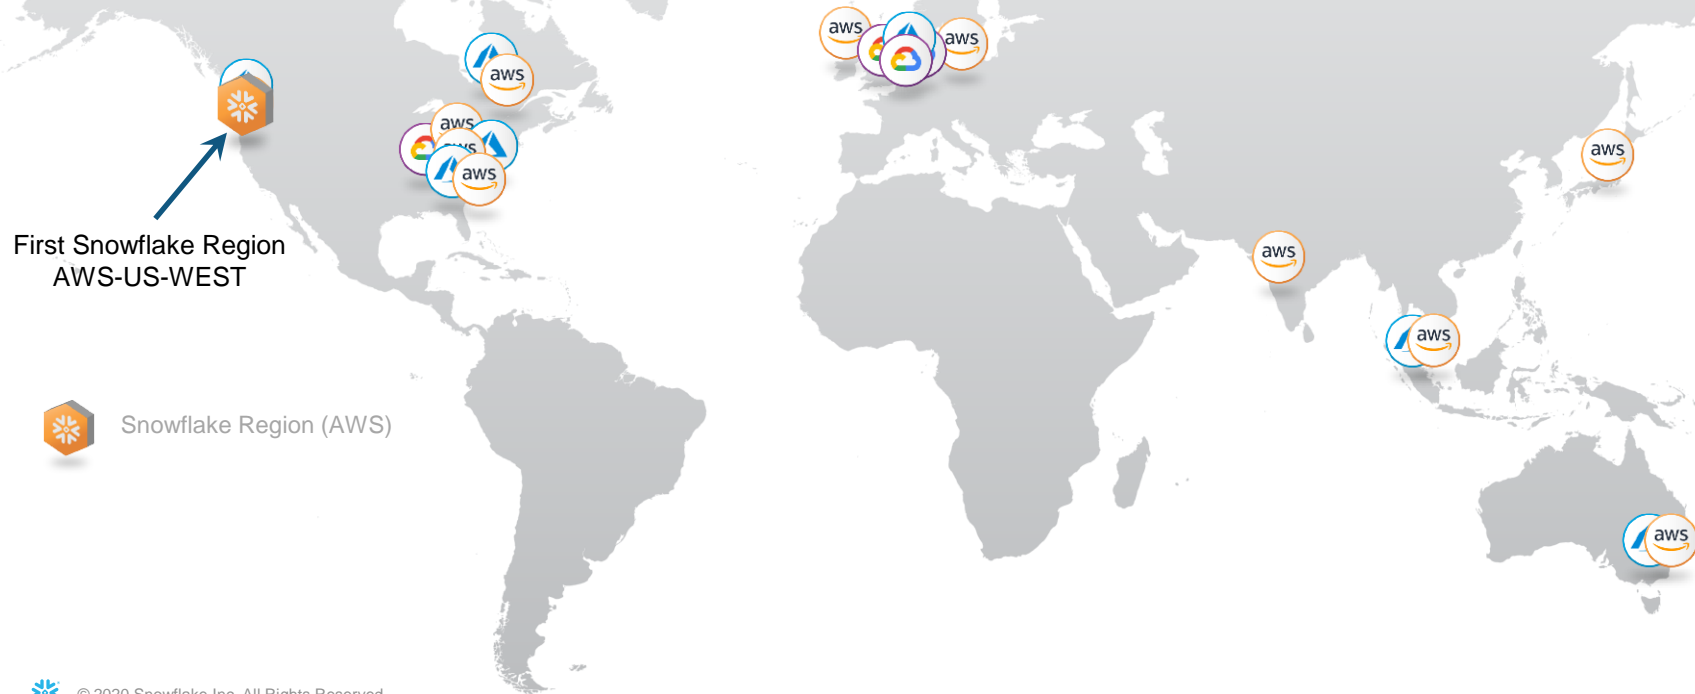


SNOWFLAKE DATA CLOUD REGION



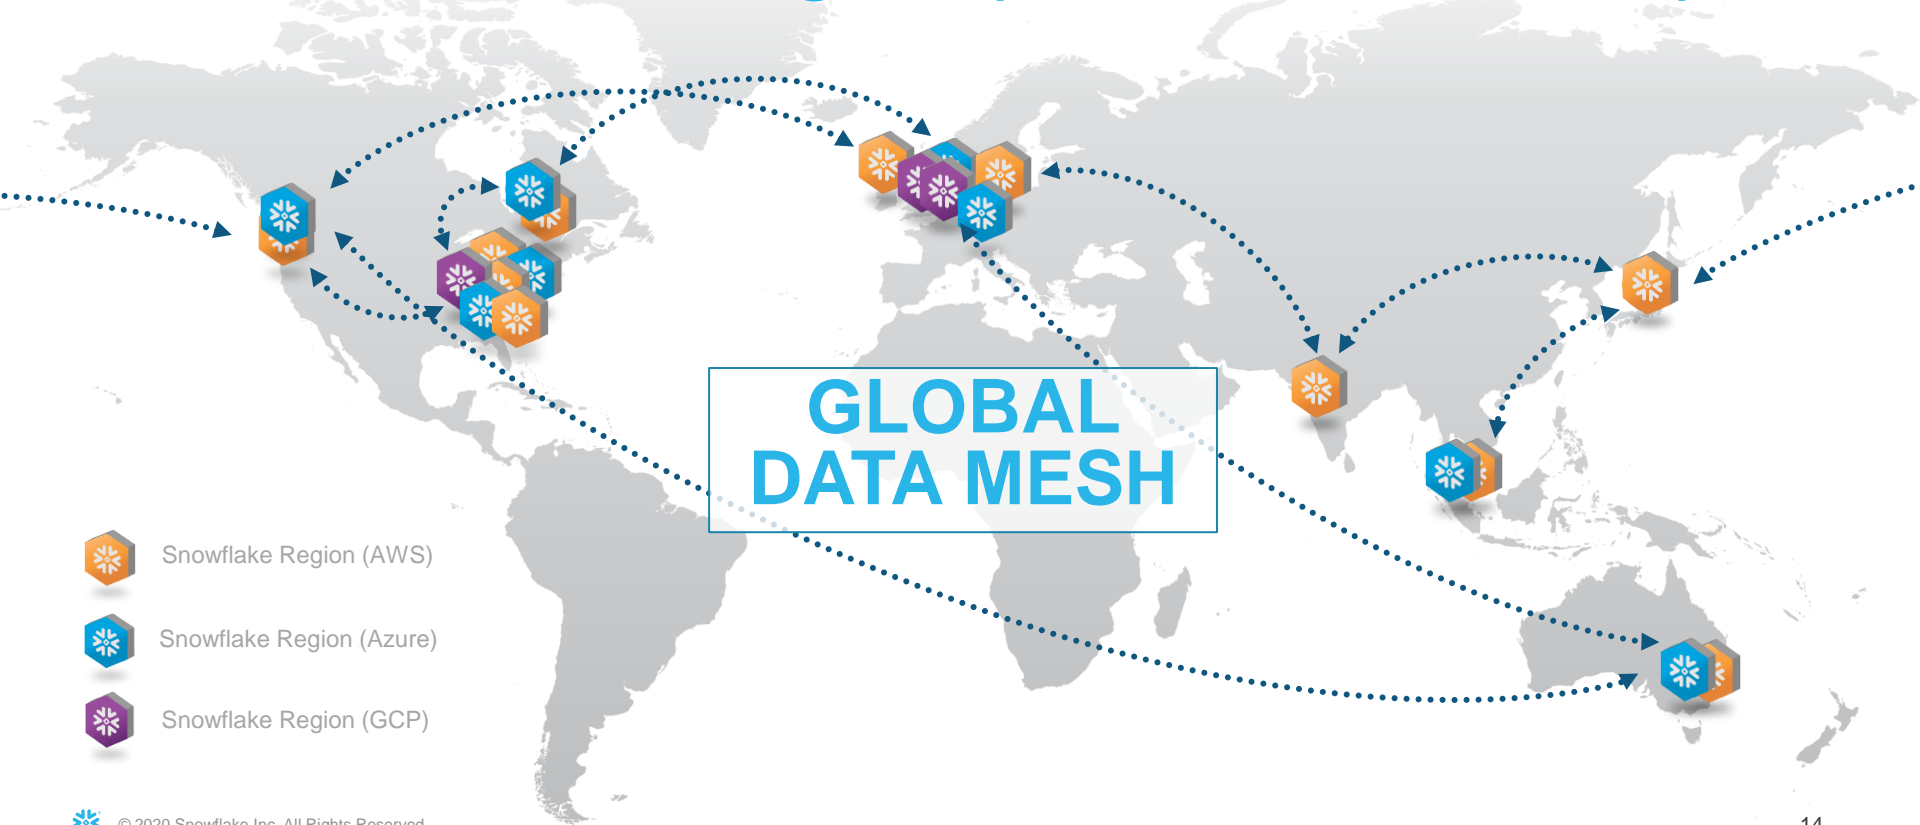
SNOWFLAKE DATA CLOUD (2015)

Single Data Cloud Region (AWS)



SNOWFLAKE DATA CLOUD (2021)

22 Data Cloud Regions (10 countries, 3 clouds)



BUILDING A LARGE-SCALE GLOBAL SERVICE

Lessons Learned

Way harder than anticipated...

- Customers expect at least 3+ 9's of availability, 24x7
- At large scale, anything will happen. Hence we need to proactively anticipate and defend
- Everything needs to be fully automated and fully adaptive
- As much as possible self-managed versus dev-ops automation
- Keeping up with exponential growth → scale cloud services and removing bottlenecks
- Weekly release without introducing (visible) regressions

... but so much faster development cycles

- We have built a top-notch and feature rich platform in only few years!
- Weekly release worldwide with single version to maintain
- Virtuous cycle – data driven development to identify and prioritize feature development
 - For example, use focus on improving DMLs and transaction processing since dominates
- Snowflake platform is extensively instrumented → we generate many terabytes of service data daily




COLLABORATIVE



DATA COLLABORATION

Traditional Way



Data providers

1. Export data to files
2. Publish schema
3. Stage files for transport

- Redundant
- Inflexible
- Inefficient
- Insecure
- Expensive

Data customers

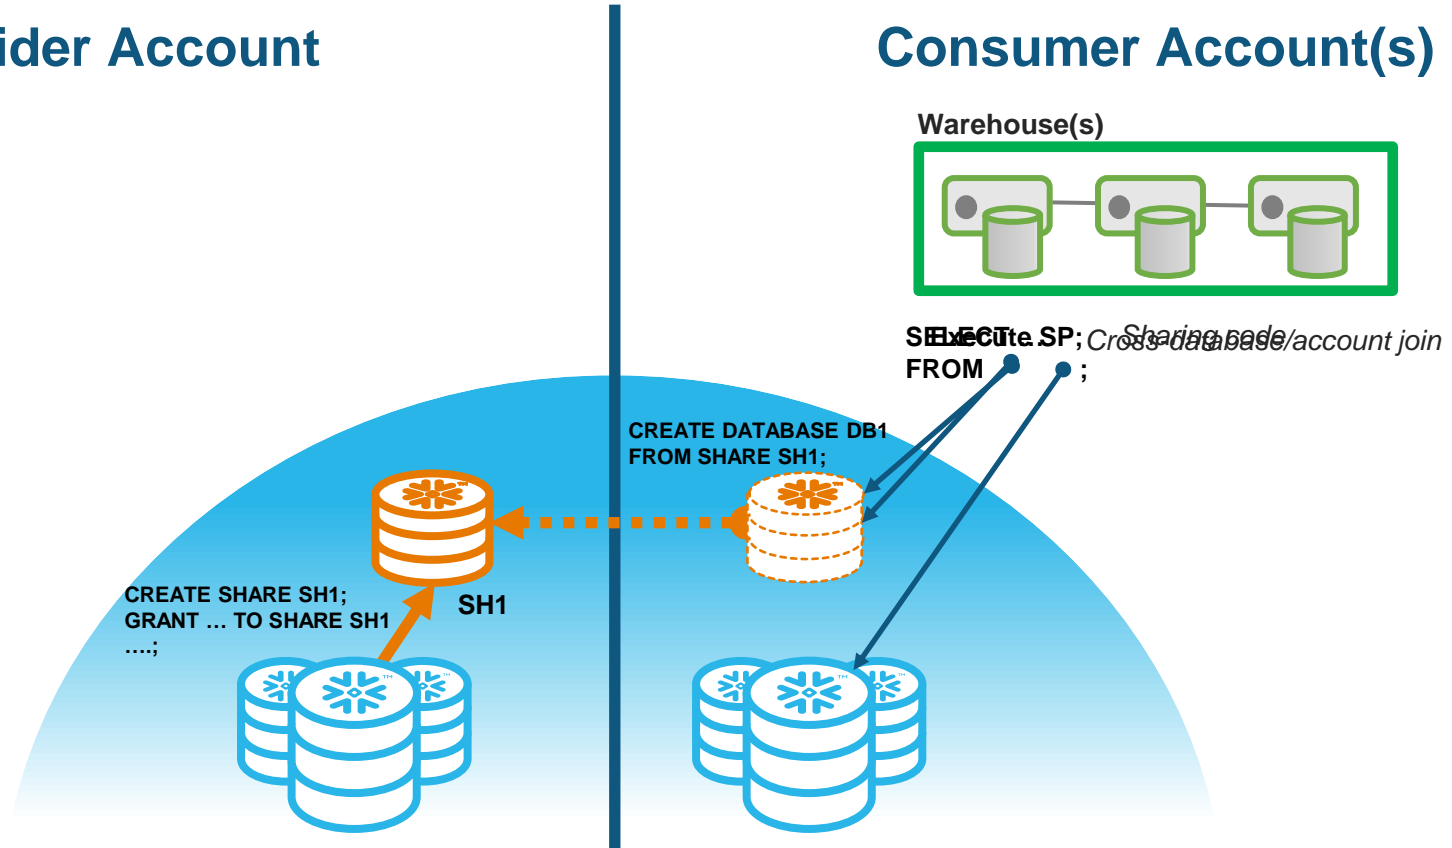
1. Additional infrastructure
2. Forced to recreate data structure
3. Delayed updates to data



SNOWFLAKE DATABASE SHARING

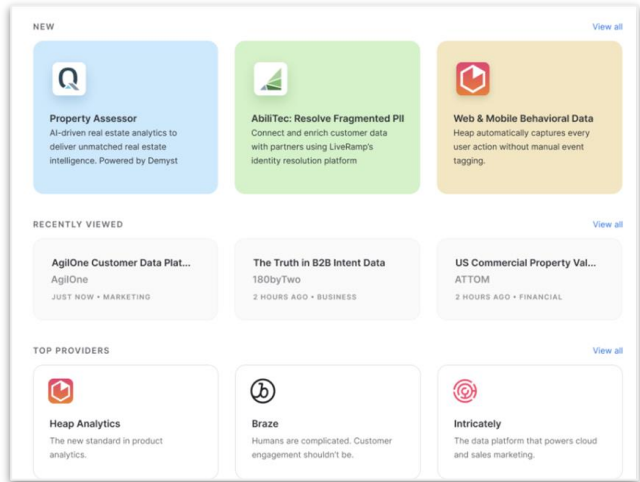
Provider Account

Consumer Account(s)



SNOWFLAKE DATA MARKETPLACE

READY TO USE DATABASES FROM MULTIPLE PROVIDERS



Live, ready-to-query data; no copying or moving



Only data marketplace with personalized data



Globally available, across clouds



Financial



Marketing



Demographic



Macroeconomic



Government




Healthcare



Business

CONNECTED THROUGH DATA

Covid-19 Database
(from Starschema)



SNOWFLAKE DATABASE SHARING

Conclusion


Secure

Live

Frictionless

Personalized

Global



SIMPLICITY



WHY SIMPLICITY MATTERS

Manage Data, Not Infrastructure!



Infrastructure

- Initial Setup
- Upgrading
- Patching
- Capacity Planning
- Storage
- Security



Physical Design

- Partitioning
- Indexing
- Ordering
- Vacuums



Data Collaboration

- Loading
- Moving
- Transforming
- Copying
- Securing



Query Tuning

- Statistic Collection
- Memory Management
- Parallelism
- Query Plan Hinting
- Workload Management



Availability

- Setup High availability
- Handle Hardware Faults
- Manage Backups



SNOWFLAKE CLOUD DATA PLATFORM

Minimal Administration



Infrastructure



Physical Design



Data Collaboration



Query Tuning



Availability

Simply load/share data and run queries

Initial Setup

Upgrading

Patching

Capacity Planning

Storage

Security

Partitioning

Indexing

Ordering

Vacuuming

Loading

Unloading

Transforming

Copying

Securing

Statistic Collection

Memory Management

Parallelism

Query Plan Hinting

Workload Management

Replication

Backups






Re-Clustering

Account Management



Infrastructure







Initial Setup		NA - Service is always on
Upgrading Patching		Automatic – Performed weekly by Cloud Services
Capacity Planning		NA - Just-in-time compute
Storage		NA - unlimited storage, spill to blob storage
Security		Automatic – Encryption, Monitoring, ...




Physical Design



Partitioning		NA – automatic at load time
Ordering		Default/Automatic Clustering
Indexing		Search optimization service
Vacuuming		NA – Immutable partitions

Data Collaboration



Query Tuning



Statistic Collection	→	Automatic – at DML time
Memory Management	→	Automatic – Cooperative Memory Brokering
Parallelism	→	Automatic – Adaptive
Query Plan Hinting	→	Robust adaptive execution strategy dynamic join filters, adaptive push down and distribution methods, join skew resilience
Workload Management	→	Virtual warehouse per workload Auto-scale multi-cluster warehouse



Availability



Setup High Availability



Out-of-box: Snowflake Architecture Multi-AZ
Disaster Recovery: Cross-region Replication

Handle Hardware Faults



Automatic: Snowflake Cloud Services detects
and replace faulty hardware

Backups



Automatic: blob storage with 11 9's durability,
Undrop, Clone as-of, time travel, Fail-Safe



CONCLUSION



SNOWFLAKE DATA CLOUD

Worldwide Web of Data



Single System

Simplicity

Collaborative





THANK YOU

