

# The state of databases in the current landscape



July, 2023

Michael Widenius

Creator of MySQL and MariaDB

CTO @ MariaDB

# Overview

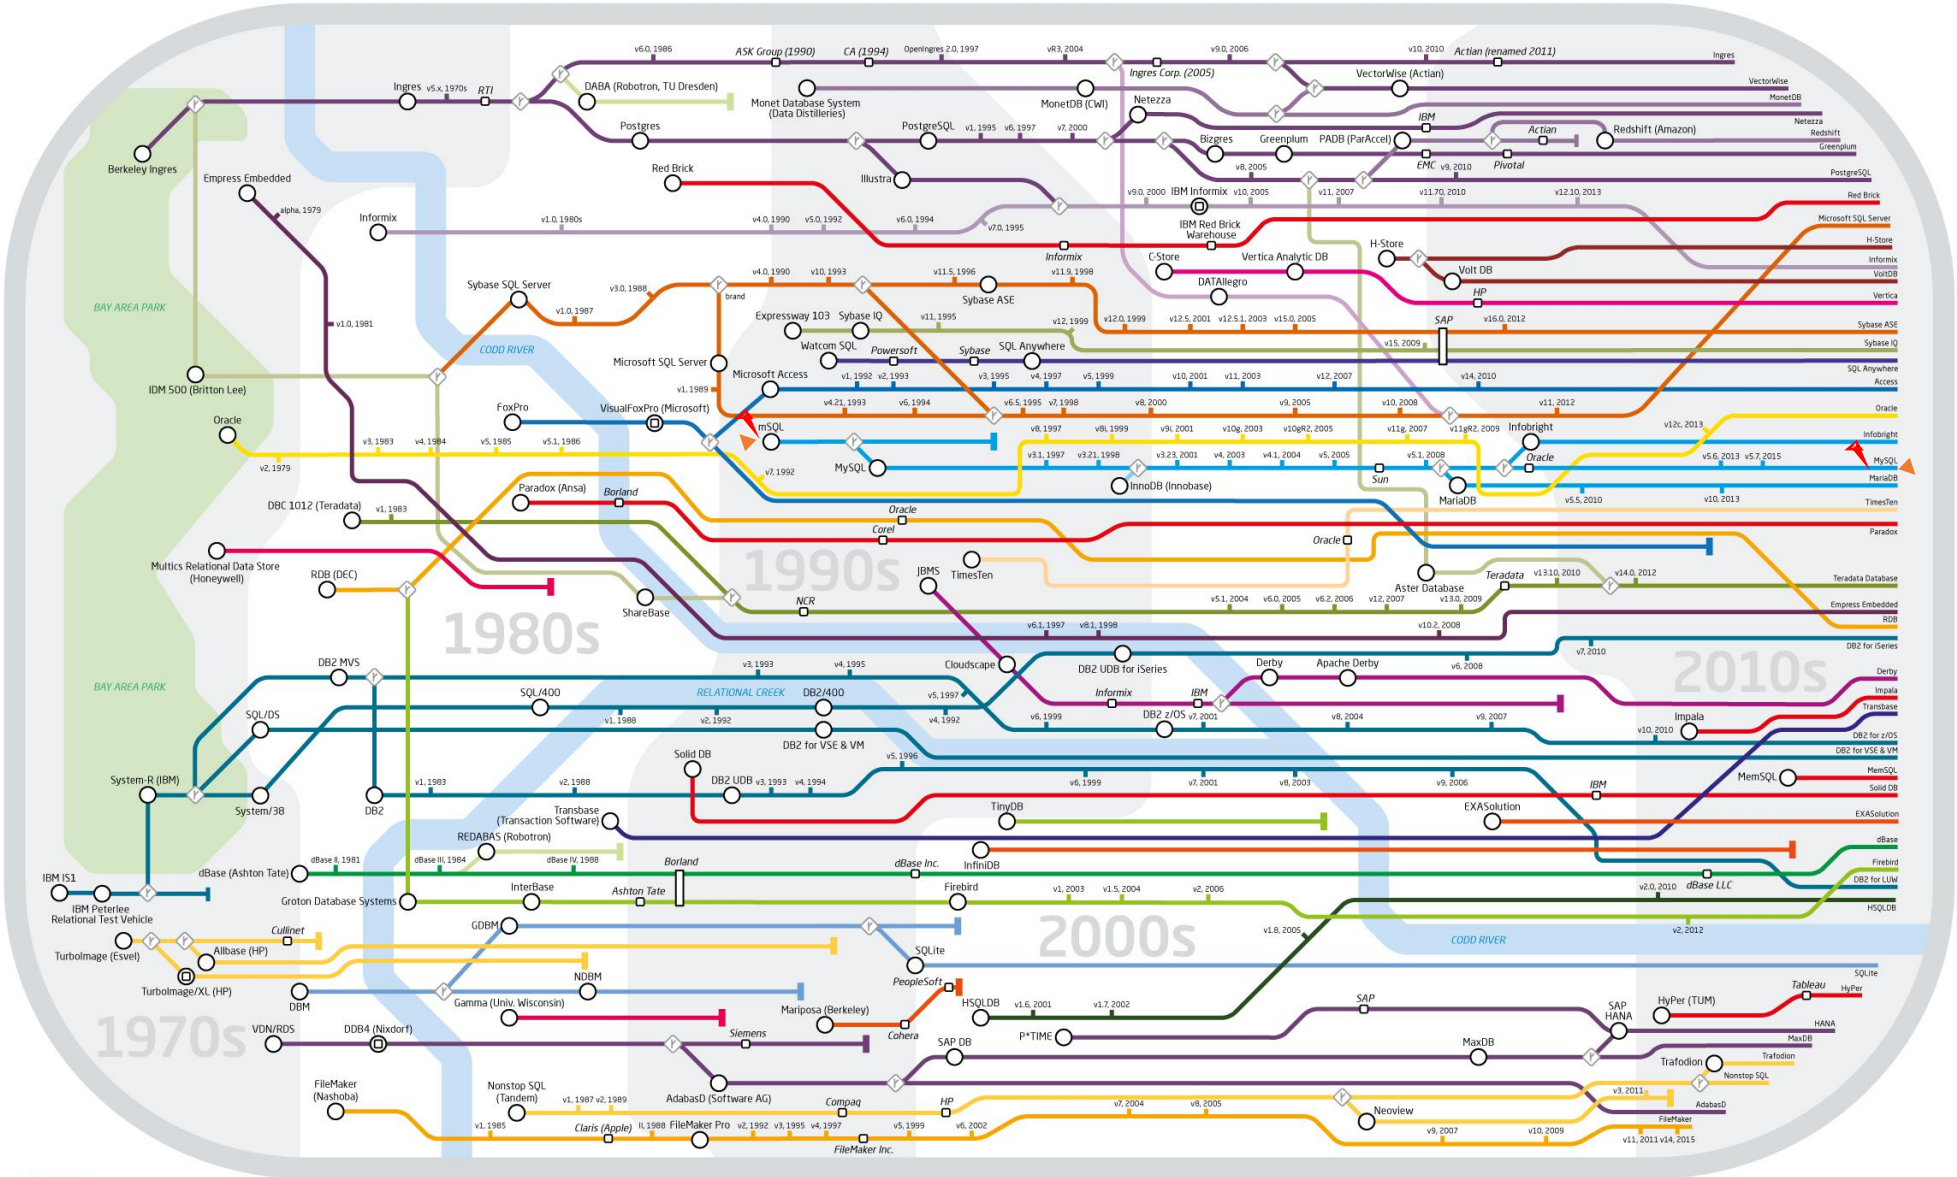
- The History of RDBMS and MariaDB
- Big Data is driven by Open Source
- AI and databases
- Databases and the cloud
- What I have been working on lately

# The history of RDBMS

# Genealogy of Relational Database Management Systems

We  
Start  
Here

Here  
Now



**Key to lines and symbols**

- DBMS name (Company)
- Acquisition
- Versions
- ⊥ Discontinued
- ◇ Branch (intellectual and/or code)
- Crossing lines have no special semantics

# What made MySQL successful?

- When MySQL was started (1994) there were big players such as Microsoft, IBM, Oracle, all having a significant portion of the market.
  - Internet was new and everyone needed a web-optimized DB
- However, the big players did not see the Internet as a viable business platform!
- MySQL was already proven stable before release (used for data warehousing and web)
- We created a “Virtual company”, which made it easy to find good people
- New “free” license scheme (this was before Open Source)
  - Free for most, a few have to pay
  - Second program (ghostscript was first) to use dual licensing.
    - MySQL was first to do it with GPL.
- Very easy to install and use (15 minute rule)
- Released source and tested binaries for most platforms
- MySQL was a needed, stable and easy to use product with the right price

# What made MySQL successful?

- MySQL started from the bottom–up, providing a cheap (or free) solution for the web industry.
  - We were friendly and helpful towards community
- I personally wrote 30,000+ emails the first 5 years to help people using MySQL
  - MySQL then started growing in other industries (enterprise sector).
  - MySQL followed an open source development model.
- "1 customer in 1000 actually pays" – Still enough to grow very fast.
- The community provided testing, marketing and simple support needs.
- MySQL Ab provided full support for those that wanted/needed it.
- We (the MySQL founders) waited with investments until product was “good enough”

# Why MariaDB was created

“Save the People, Save the Product”

- To keep the MySQL talent together
  - To ensure that a free version of MySQL always exists
  - To get one community developed and maintained branch
  - Work with other MySQL forks/branches to share knowhow and code
- 
- After Oracle announced it wanting to buy Sun & MySQL this got to be even more important.

# How was MariaDB disruptive?

MariaDB follows in the same original footsteps of MySQL (which Oracle has not done):

- It is a true to Open Source project, following proper Free Software / Open Source practices.
- Development happens in the open, working together with the community.
- **MariaDB Foundation** was created to ensure that MariaDB would always be Open Source.
- This process made MariaDB stand out.
  - MariaDB has been integrated into most major Linux and other free OS distributions as the default “MySQL” variant.



# MariaDB Ecosystem

## **MariaDB Foundation** **Works with the community**

Builds and tests binaries  
- Develops MariaDB buildbot

Drives adoption  
- Works with OS to ensure MariaDB is included everywhere

Works with community developers  
- Reviews architecture and patches  
- Approves and pushes changes

Insures that MariaDB is always free.  
Founded through sponsorship's.

## **MariaDB Corporation** **Works with customers**

Provides paid support and subscriptions for MariaDB (Enterprise and Community)

Employs most of the MariaDB developers  
- Main driver of MariaDB development

MariaDB Enterprise  
- Longer End-of-life  
- Stable features are backported to earlier versions to minimize needs for upgrades.

Provides NRE (paid development of new MariaDB features).

# MariaDB future plans

Monty's view

- Work with customers to help them migrate from commercial closed source database to MariaDB
  - The MariaDB Oracle compatibility layer makes this easy
- Improve optimizer to handle very complex queries
  - This is important when moving complex application to MariaDB
  - MariaDB 11.0 has a new cost model that greatly improves handling of complex queries.
- Work closer with SAS database providers to make MariaDB work better in their environment
  - MariaDB multi-tenancy feature is part of this collaboration.

Big data is driven by  
Open Source

# Big data driven by Open Source

- MySQL / MariaDB
  - SQL database with flexible replication, ColumnStore and Spider for scale out
- Apache Cassandra
  - Tunable consistency
  - Keys map to multiple values, which are grouped into column families
  - CQL language
- MongoDB
  - Stores structured data as JSON like documents
- CouchDB
  - MVCC, ACID, eventually consistent
- Hbase (Hadoop + distributed file system)
  - Written in Java. Compression, in-memory operation and Bloom filters (Is data part of a set).
- HAWQ, Pivotal HDB
  - Hadoop + SQL
- Redis
  - In memory database (+ snapshots to disk)
  - Optional durability
- Clickhouse
  - Analytical column oriented database

# Why use NoSQL

- Faster replication
- Fast and easy key/value access
- Data is often stored in memory
- Note that with similar memory resources you can usually keep SQL data in memory too.
- Can handle unstructured data
  
- In traditional SQL, one can't easily implement a web store
  
- Most NoSQL vendors are considering adding SQL support.
  - Like Atlas SQL for MongoDB (not open source)

# Why NOT to use NoSQL

- Lot of data duplication
  - Relational database are designed to keep only one copy of a relation (like a customers data)
- Can't (easily) join/combine data
- No standardized language; Hard to move data and applications between systems
- Complete lock in to one system
- Normally few connections to computer languages or existing applications that need a database.
- Allows one to initially ignore solving the database layout problem as one can store data 'as such'
- Initially easier to use, MUCH harder later
- No sustainable scalable business models for most open source NoSQL solutions (as most of them rely on the BSD or Apache licenses)

# Problems to overcome with BIG data

- Storage (not normally the big problem anymore)
  - Memory (Very expensive)
  - SSD (Expensive but slower)
  - Hard disk (Inexpensive. Slow when doing random reads, fast on sequential reads. The "tapes of tomorrow")
- Access patterns
  - Some access patterns are good for big data (scanning in parallel) while others are hard
    - Even simple joins can take days on petabyte tables
- Get it up when things fail
  - Recovery; Can take days or weeks
  - Get the caches warm (you need > 90 % hit rate with buffers!)
- Replication and hot standby
  - Must have, but makes big data more expensive

# Definition of big data is changing

- When I worked with data warehousing (1986–1993) , big data was (all credit card transactions for one Sweden's oil companies):
- 1M users, 4M transactions / month (30 byte / transaction)  
10 years of data = 1.3G
- Handled by a Sun SparcStation, 25Mz, 32M memory. 2G hard disk
- CPU's are now 700 times faster and have 10,000 times more memory.  
SSD/Hard disk seek speed is 6000 times faster than in 1986
- While machines have got faster, most data is still small (except for social and behavior data) and can easily be handled by one machine.



# Why big data solutions are hype

- Most companies will not have as much real data as Facebook, Twitter, Bilibili, TikTok, etc. that needs to be accessed "at once"
- Solutions that they have to use are not applicable for others.
  - In the near future one can run what most companies think is big data on a few machines:
  - Memory now costs about 2100\$/T; Most can afford 200G of RAM.
  - SSD are now 50\$/T (Read: 560–3100MB/sec 98K IOPS, Write: 520–2250MB/sec write).
  - You can get a 100Tb SSD from Nimbus and a 960Tb from Dell. 16T can easily get bought.
- MariaDB today can easily handle 0.5T of active data. It is only after 1T when one needs to consider analytical databases.
- There are very few users that need more than 0.5T of data.
- The most common database size for big enterprises are **100G** for their largest production databases.
- 90% of queries processes less than **100 MB** of data.
  - See <https://motherduck.com/blog/big-data-is-dead/> for details

# AI and Databases

# AI and databases

- Personally I don't see AI helping optimizing queries inside the database (data is changing constantly and the optimizer is already using statistics and histograms to make decisions).
- Optimizer also has to be very fast and handle a lot of concurrency with a limited amount of memory.
- AI can be used to optimize SQL queries for applications.
- AI can be used to translate natural language to SQL, which is useful for interactive queries but too slow to be used in production for applications.
- When it comes to programming, I would only trust AI to:
  - Initially set things up for a common problem (write a script that does 'easy explanation')
  - Try to find 'simple mistakes/bugs' in code like possible buffer overflows, missing arguments, things that can easily be simplified etc.
  - Translating things from one language to another.
  - Do a simple change that can be repeated in many parts of the code.
    - Add a parameter of this type to this function and ensure that all callers are fixed.
    - The resulting code has to be carefully reviewed!

# AI and databases

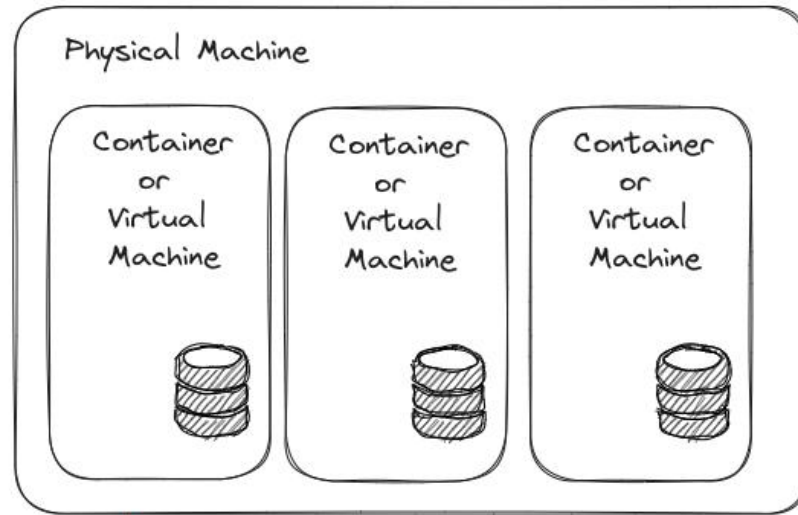
- I would never even try to use AI for doing changes in a complex project, like adding catalogs to MariaDB.
  - There are no similar patterns in existing software that it could be trained on.
  - Too many things in MariaDB are interconnecting in not obvious ways (if one does not know the code)
- If the AI would produce 'hard to understand code', I could not use AI to explain why it did things that way.
  - All MariaDB developers I am working with are experts and can be put on solving very complex problem in their domain. I don't see anyone being replaced by AI during their or their children's lifetimes.

# Databases and the Cloud

# Database scaling for SaaS providers

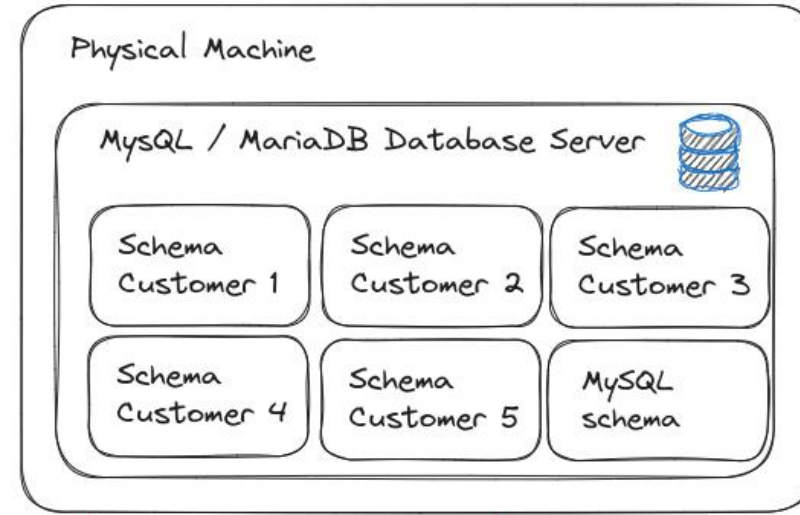
- SaaS providers want to optimize the number of customers they can host on their hardware.
- When offering managed services on MySQL/MariaDB, SaaS providers have the following options:
  - Create separate VMs for each customer.
    - Large overhead – costly
    - A typical “idle” MariaDB Server requires ~1GB of memory
    - Best user experience
  - Use shared instance
    - Force user restrictions, only grant limited database access.
      - cPanel has this model, many other database management systems share it
    - Very limited overhead
    - Poor user experience, artificial limitations enforced on users
    - Affected by “noisy neighbour”, hard to track

## Container / VM Approach



→ 1 GB minimum per VM

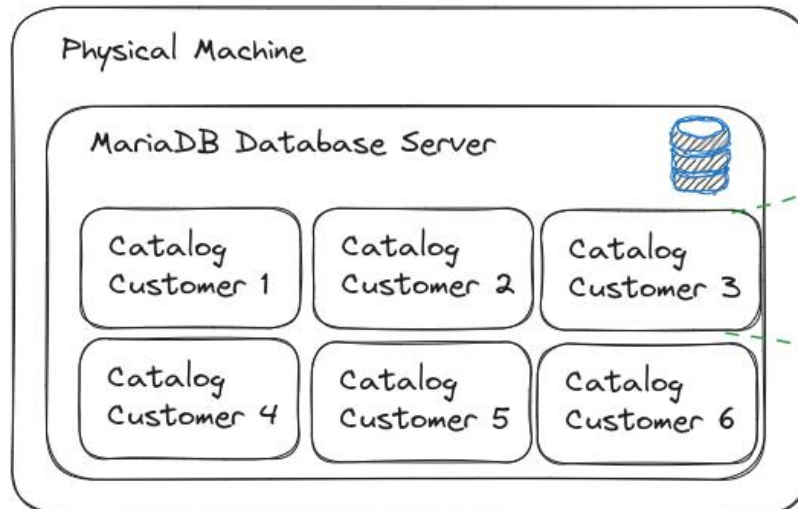
## Shared Server / Separate Schema



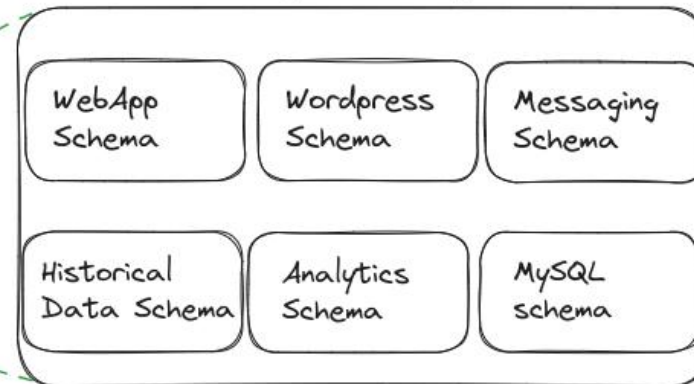
→ Customers have limited schemas

Catalogs

## Catalogs



## Catalog For a Single Customer



# New catalogs feature for MariaDB

- Catalogs provide the best of both worlds.
  - Shared instance
  - A catalog looks like a normal MariaDB Server
  - No user limitations
    - Each user has full control over their catalog
    - Root access on the catalog
- Catalogs still have the problem of noisy neighbour
  - However, statistics are now collected for each catalog
  - So it's trivial to detect the problematic catalog
  - Move the problematic user to a bigger machine, increase costs
- Catalogs can also enforce quotas
  - Each catalog have their own configuration file with their own limits
  - Force users to stick to certain performance limitations.
- Sysadmins have control over all catalogs.



# Catalogs feature

- Catalogs enable hyper-scaling for SaaS providers.
  - For basic users, one can now host up-to 100x more users on the same machines.
- In MariaDB, this feature is still under development
  - If you would like to steer the roadmap for this, now is the time!
    - Tooling
    - Performance optimizations
    - Specific functionality

# Career path for programmers

# What I have been working on lately

I am still actively coding!

- Improving the MariaDB optimizer to handle complex queries for big databases
  - This is based on input from MariaDB customers and users
  - I spent almost one year on redoing and improving the cost model of the MariaDB optimizer for MariaDB 11.0.
    - This was what I was working on during my 10 days of Covid quarantine during my last trip to China!
  - Initial test shows that most of the recent queries that has caused problems with MariaDB 10.6 optimizer works very good in MariaDB 11.0!
- The last 6 months I have been working on the multi-tenancy catalogs for MariaDB.
  - This was based on input from several SAS provider at the last 2 CloudFest conferences.

# Career path for programmers

- I created MySQL and MariaDB. After the initial setups (where I handled everything), I have always hired people to do management, customers and leading the company so that I can continue to **focus on architecture, development**, and leading the MySQL/MariaDB developer teams.
- A very common question I get in China is what is the best career path for a developer? Should I become a manager or start doing something else?
- Good coders are hard to find!
  - It takes a long time for a coder to reach their peak (8–10) years of experience.
  - Good programmers are valuable and can produce quality code for LONG time (>> 70 years)
- Don't waste a good coder by “promoting” him to a manager.
  - A good coder is not necessarily a good manager.
  - Promote them by giving them more responsibility:
    - Architecture – design bigger and more complex systems
    - Help others to work on their code (but not manage!)
- My advice to managers:
  - **Do not** follow the “**Peter Principle**” to promote people until they reach their level of **incompetence**!
  - The expected salary raise/year in China is not suitable for long term programmers (35 year rule!)



**Thank you**